



Building Streaming Application Using Kafka

Kodjo Klouvi | kodjo.klouvi@gmail.com

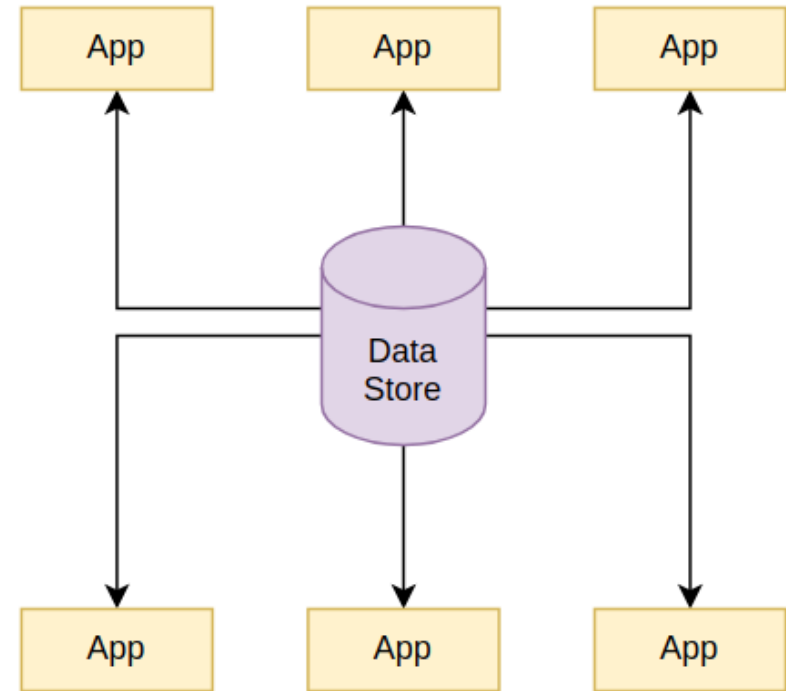


What is Kafka?

- Event-driven architecture: publisher – subscriber topology
- Not a client-server architecture:
 - multiple brokers playing the same role
 - any broker can be used to target the cluster (bootstrap server for producer/consumer)
- Used to build streaming applications and data pipeline
- High scalability, fault tolerant, support data retention
- Many options to configure and some require advanced expertise!

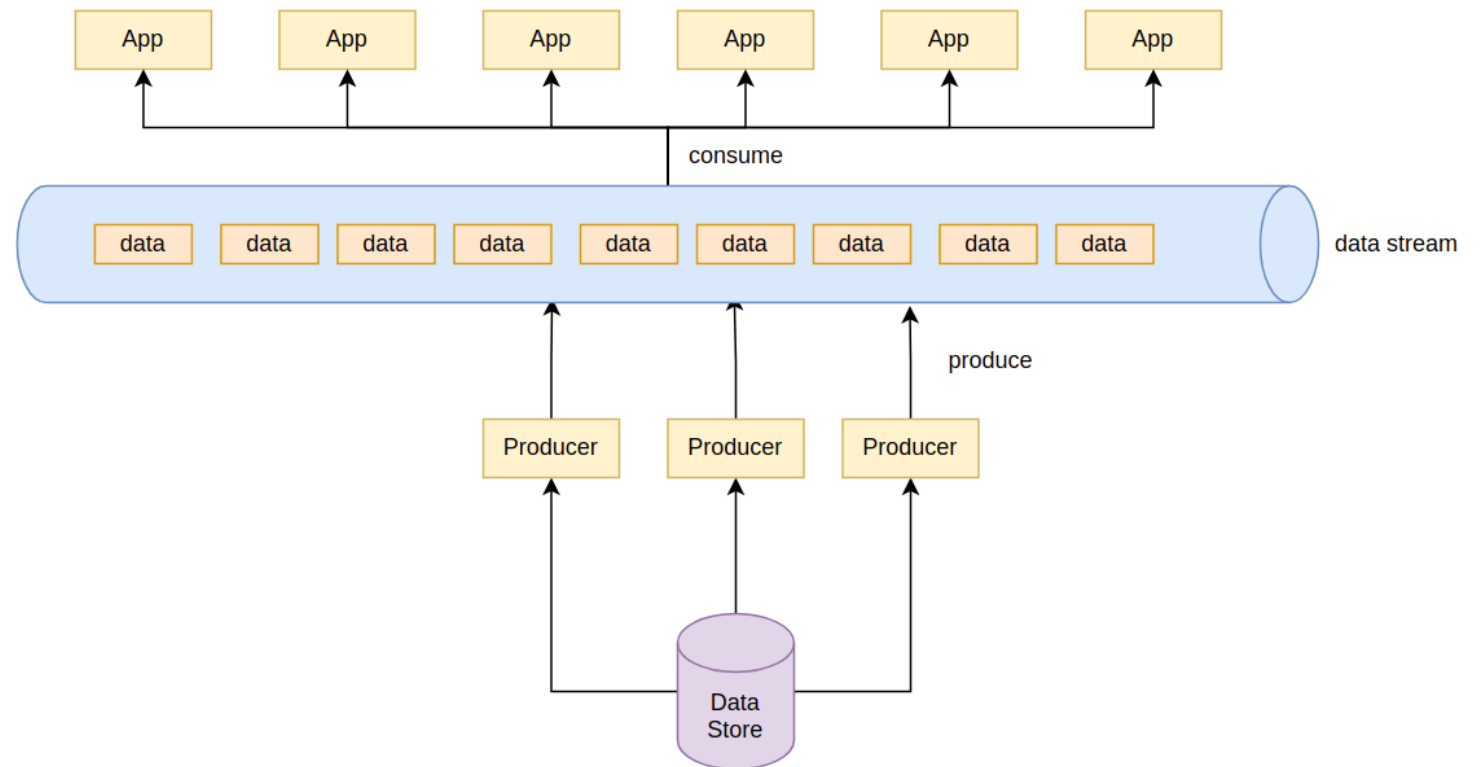
Flashback...

- Multiple applications access the same data at the same time and apply the same transformations/operations
 - Risk of bottleneck
 - Massive resources consumption
 - Duplication
 - Heavy to maintain
- Synchronous actions (read - write)
 - Time consuming
 - Not parallelable



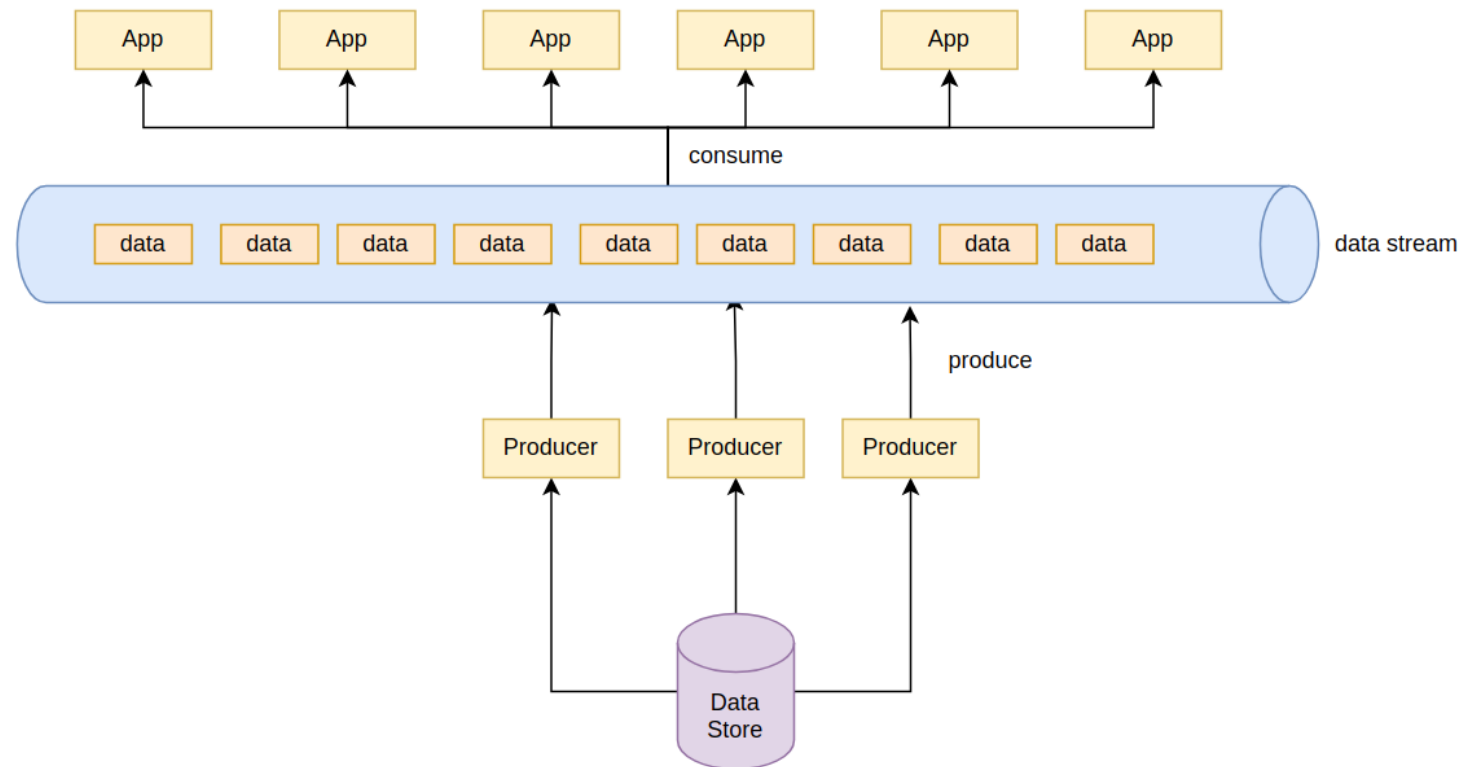
Kafka Architecture

- **Message queue:** data structure to store messages (events)
- **Producer:** application that publishes messages into the queue
- **Consumer:** application that reads messages from the queue



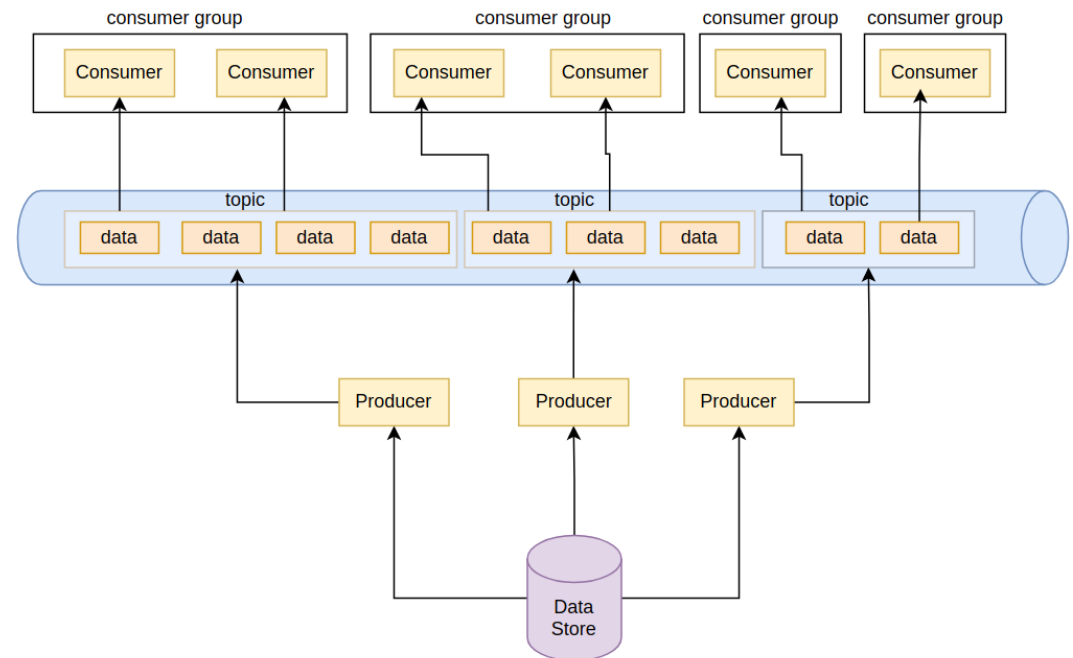
Kafka Architecture

- **Broker:** application that manages the data.
- **Zookeeper:** manage the infrastructure, the connectivity (network discovery) and meta data
- **Asynchronous actions**



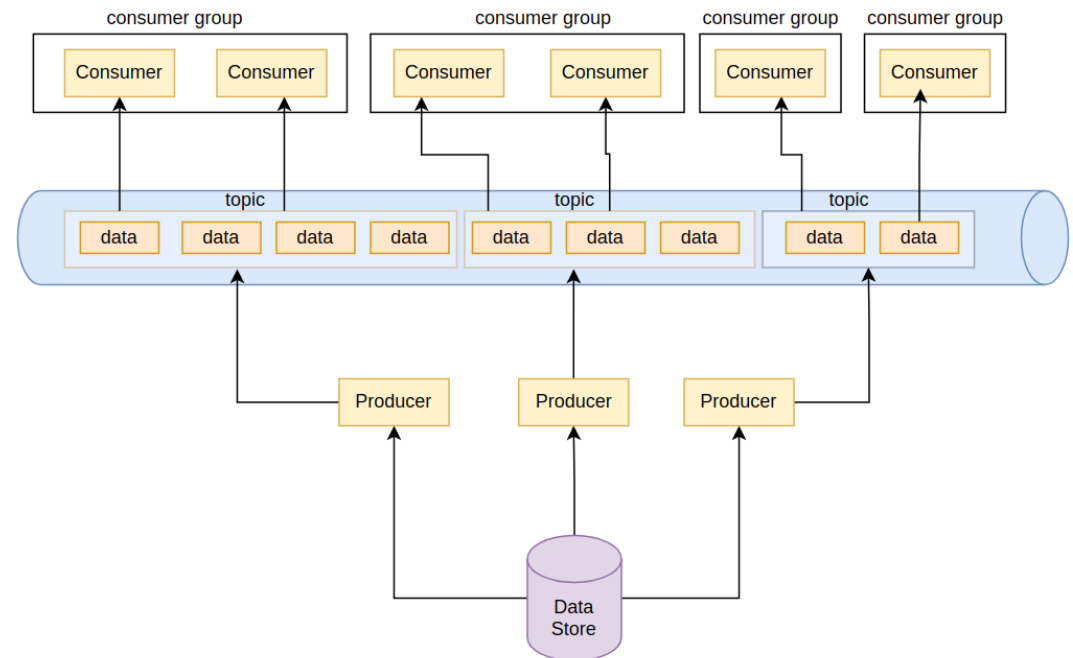
Kafka Architecture

- **Topic:** data structure to store the same type of data (for the same purpose)
- **Partition:** division of a topic. Each topic can have multiple partitions allowing horizontal scalability across many servers.
- **Offset:** unique identifier assigned to each message within partition. The position of the message in the partition.



Kafka Architecture

- **Consumer group:** group of applications/consumers that read the same topic – they work in parallel using partitions and offset paradigm
- **Multiple brokers:** data are stored on multiple nodes (replication)
- **Replication:** each message is stored on multiple nodes (replication factor) and has its own leader (broker)
 - fast and parallel read
 - fault tolerant, etc.





Kafka Architecture

- Only one broker per machine
 - Kafka doesn't support horizontal scalability at the machine level
- Each broker has its own ID (broker.id)
- Data retention period: keep data (in the log) for some period
- Programming language agnostic
- Sensitive to the data structure/format, serialization/deserialization



Use Cases

- Real-Time Data Streaming and Analytics
- Data Pipeline Modernisation (ETL)
- Microservices Communication
- Event Sourcing
- Log Aggregation
- Stream Processing (Kafka Stream API)



Use Cases

- Data Integration Across Systems (Kafka Connect)
- IoT Data Management
- Replication and Disaster Recovery
- Edge Computing



Hands-On

- <https://github.com/osekoo/hands-on-kafka>
- Preparing development environment
- Practice 1: Get started
- Practice 2: Implement online word definition search (dictionary)
- Practice 3: Homework (advanced dictionary)



PyCharm

- Python Integrated Development Environment
- Multiple modules available for Python
- Enable debugging, Python applications packaging, etc.
- Even more...

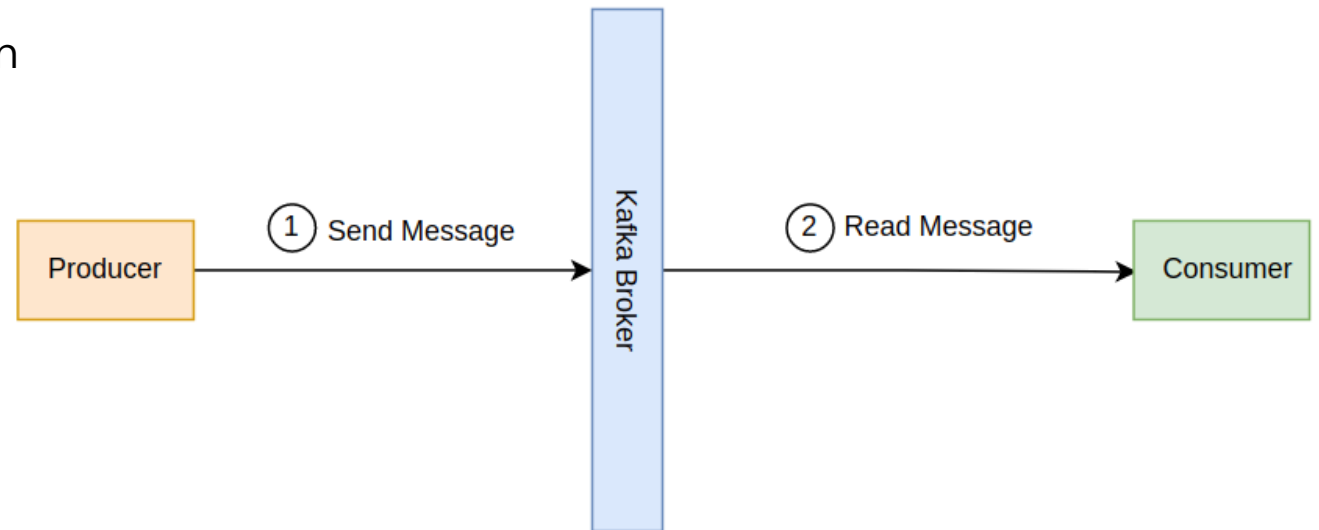


Docker

- Platform as a service (PaaS)
- OS abstraction
- Self-contained applications running in containers
- Simplify applications delivery/distribution
- Enable horizontal scaling

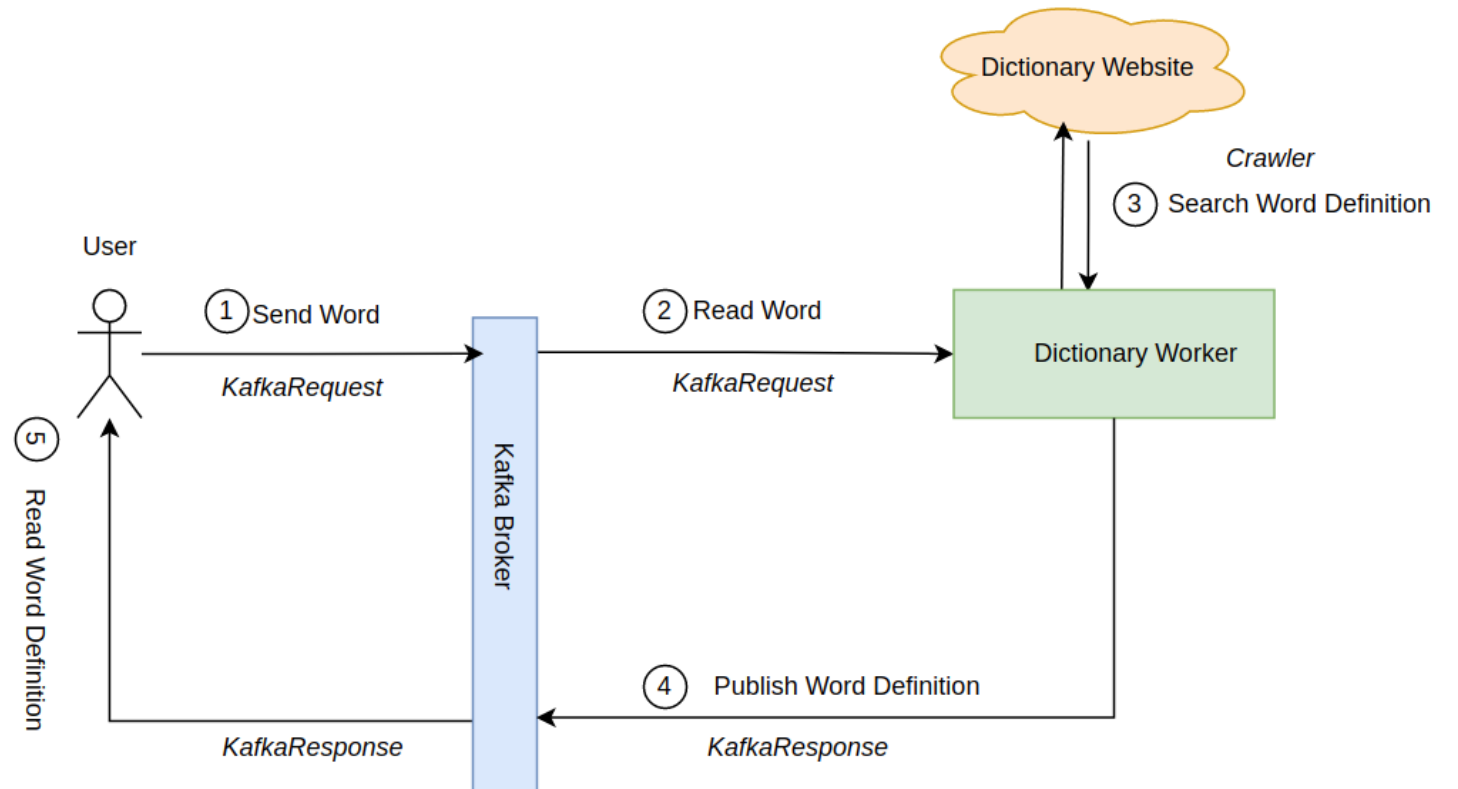
Practice 1: Get started

- Build a very simple Kafka based application
- Implement "*push-pull*" message queue



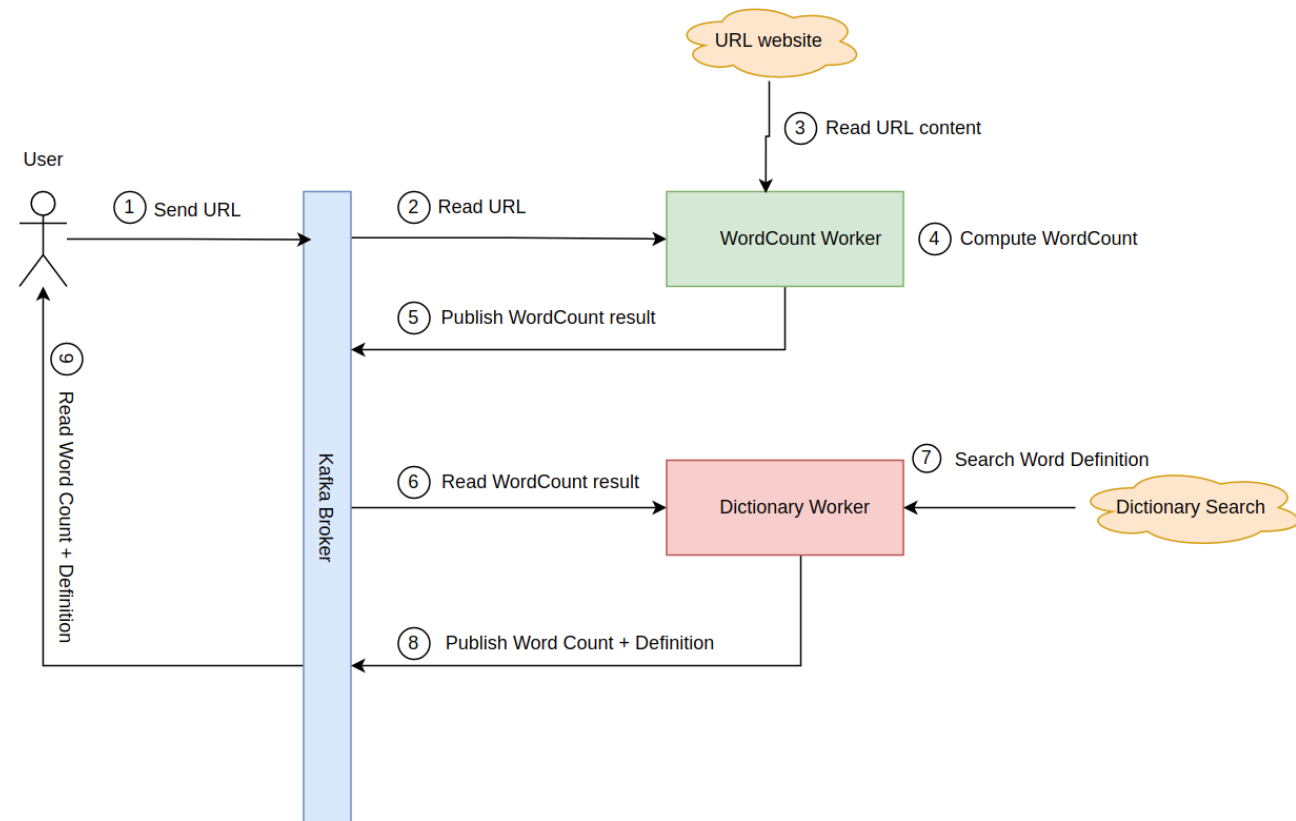
Practice 2: Online dictionary

- Build a Kafka based word definition search dictionary
- User specifies his nickname and the dictionary language (FR, EN)
- He/she interacts with the application by sending a word
- Our application search for the word definition by scrapping online dictionaries



Practice 3: Advanced dictionary

- Homework
- Searching word definition for website
- Extend language support (ES, CN, GE, etc.)
- Further details are available on Github





Troubleshooting

- Raise any question during the session or via email kodjo.klouvi@gmail.com